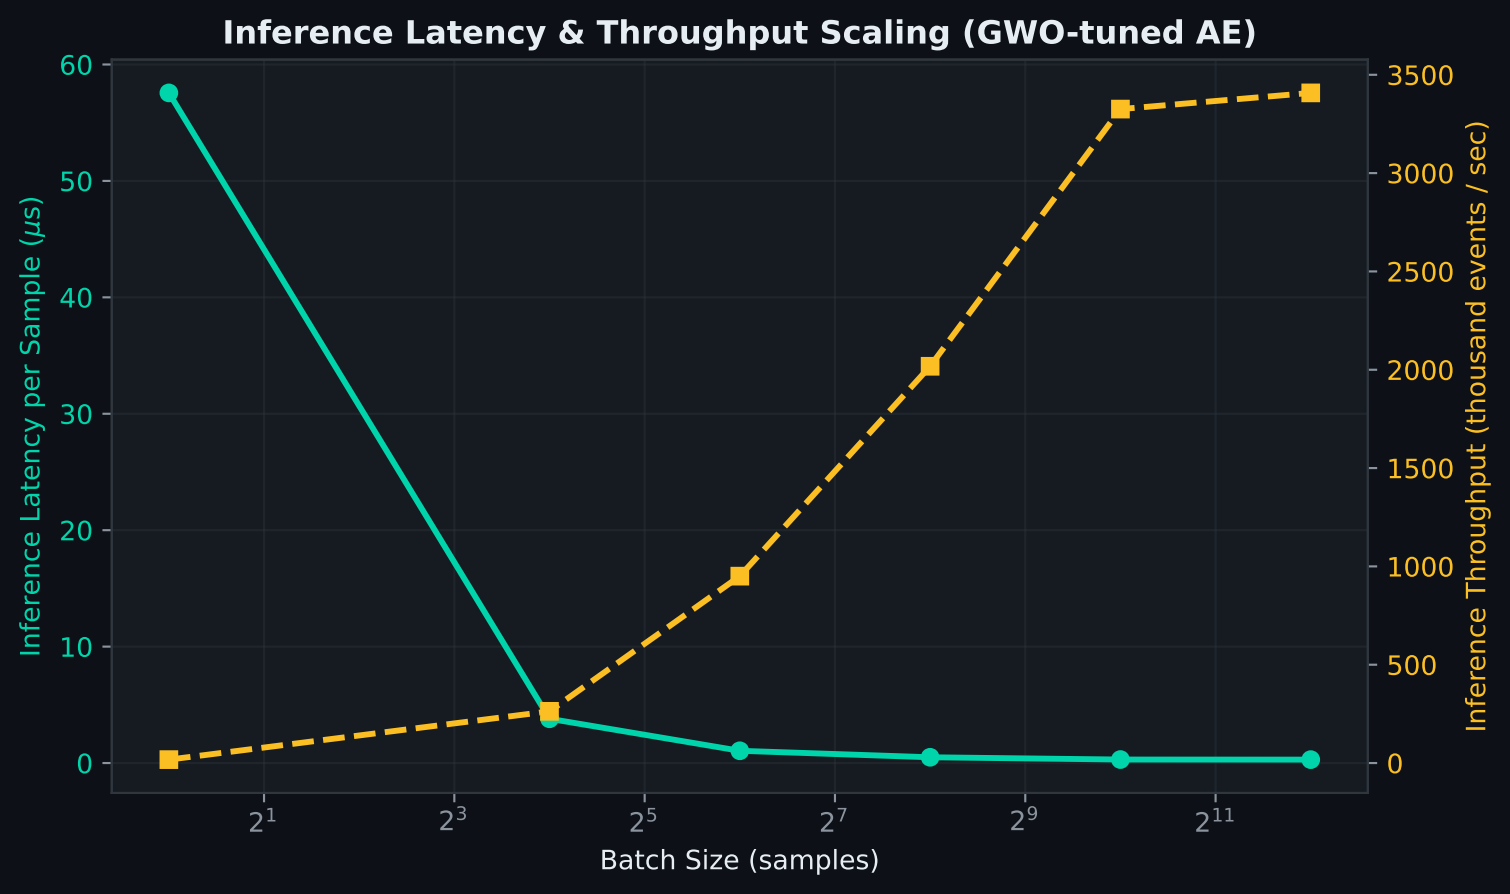
Inference Latency & Throughput Scaling (GWO-tuned AE)
Inference Latency per Sample (μs)
Inference Throughput (thousand events / sec)
Batch Size (samples)
0
10
20
30
40
50
60
0
500
1000
1500
2000
2500
3000
3500
2^1
2^3
2^5
2^7
2^9
2^{11}